

SAHIL KIRAD

+91 80804 01915

✉ kiradsahil882@gmail.com

🌐 [linkedin.com/in/sahil-kirad](https://www.linkedin.com/in/sahil-kirad)

🐙 github.com/sahilkirad



sahilkirad.vercel.app

Experience

Arealis, Inc.

June 2025 – August 2026

Full Stack AI Intern

Pune, Maharashtra

- Owned and architected **zord-prompt-layer**, a RAG-based prompt service powering a fintech product, enabling tenant specific, context aware answering for operator queries.
- Designed the retrieval pipeline using **Pinecone** for vector search and **PostgreSQL** for structured data, with hashing based mechanisms to ensure secure, tenant isolated data access.
- Built an asynchronous, event driven backbone using **Kafka** (producer consumer architecture) integrated with **Redis** caching, **improving throughput by 40% and reducing manual operational overhead by 35%** while maintaining **99.9% uptime** across the prompt layer.
- Represented Arealis at E-Summit 2025 (IIT Bombay Startup Expo), India Insurtech Summit 2026, demoing AI solutions to **50+ industry professionals** and generating **15+ high-intent B2B leads**.

Projects

CodeGuard: AI Code Review System | Qwen2.5, QLoRA, LangGraph, RAG, FastAPI

2025

- Developed an automated PR analysis system using **QLoRA** fine-tuning, **improving bug detection accuracy by 22% and reducing false positives by 18%** over baseline models.
- Integrated RAG with Pinecone and leveraged LangGraph for asynchronous processing of large-scale repositories.

VeritasAI: Agentic Investment Platform | Vertex AI, Cloud Run, Pub/Sub, Next.js

2025

- Built an agentic due diligence platform **cutting startup evaluation time from 48 hours to 15 minutes** using Gemini-based workflows.
- Optimized infrastructure with Cloud Run and Pub/Sub, **scaling to 800+ concurrent users at sub-200ms latency**.

Support Ticket Classifier & Auto-Router | FastAPI, Next.js, PostgreSQL, MLflow, Docker, AWS

2026

- Built an end-to-end MLOps system to classify support tickets by category and urgency with confidence-based auto-routing and human-in-the-loop workflows.
- Implemented retraining pipelines with MLflow tracking and feedback-driven learning for continuous model improvement.

NexusAI: Recommendation Engine | Python, FastAPI, LangChain, Redis, PostgreSQL

2025

- Architected a recommendation API with explainable reasoning via LangChain, improving simulated user engagement metrics.
- Integrated Redis caching and a business rule dashboard to reduce latency in high-traffic environments.

Publications

Published research paper titled "A Comparative Study of Parameter-Efficient Fine-Tuning Techniques for Large Language Models" in JETIR. [Paper Link]

Technical Summary

AI & LLMs: Fine-tuning (QLoRA), RAG, LangChain, LangGraph, Agentic Workflows, Vertex AI
Backend & Cloud: FastAPI, Node.js, GCP (Cloud Run, Pub/Sub), Docker, CI/CD, Microservices
Databases & Systems: PostgreSQL, Pinecone, Redis, Firestore, Distributed Systems
Frontend & Languages: React, Next.js, TypeScript, Python, C++, Bash

Education

AISSMS Institute of Information Technology

July 2023 – July 2027

B.Tech in Artificial Intelligence & Data Science (CGPA: 8.13, SGPA: 8.48)

Pune, Maharashtra

Certifications

Oracle Agentic AI Foundations Associate (2026) — Oracle University

Leadership & Community Engagement

Hackathons & Tech Conferences

Active Contributor

Global / India

- Hackathons:** 1st Runner Up, Alphabyte 3.0 (PCCOE); participated in Mumbai Hacks, Google Agentic AI, Google Gen AI Exchange, and EY Techathon.
- Open Source:** Contributed to **Google Gemini CLI** repository with PR #25313 & #26312 successfully merged.
- Cloud Native Hackathon:** Built a production-ready microservices app with CI/CD pipelines, reducing deployment time.
- Industry Engagement:** Attended L&T Finance event and API Conf 2025, gaining insights into multi-LLM systems.