

SAHIL KIRAD

+91 80804 01915

✉ kiradsahil882@gmail.com

in linkedin.com/in/sahil-kirad

github.com/sahilkirad

Experience

Arealis, Inc.

June 2025 – Present

Full Stack AI Intern

Pune, Maharashtra

- Architected a resilient multi-agent orchestration system using LLMs, reducing manual workflow intervention through DAG-based execution and fault-tolerant coordination.
- Engineered a DAG Agent and an SLA Monitoring Agent, improving pipeline throughput for real-time gateway decision-making.
- Represented Arealis at E-Summit 2025 (IIT Bombay Startup Expo), demoing AI solutions to **50+ industry professionals** and generating **15+ high-intent B2B leads**.

Projects

CodeGuard: AI Code Review System | Qwen2.5, QLoRA, LangGraph, RAG, FastAPI

2025

- Developed an automated PR analysis system using **QLoRA** fine-tuning, improving bug detection accuracy over baseline models.
- Integrated RAG with Pinecone and leveraged LangGraph for asynchronous processing of large-scale repositories.

VeritasAI: Agentic Investment Platform | Vertex AI, Cloud Run, Pub/Sub, Next.js

2025

- Built an agentic due diligence platform reducing startup evaluation time and enabling faster decision-making using Gemini-based workflows.
- Optimized infrastructure with Cloud Run and Pub/Sub to handle concurrent requests with low latency.

Support Ticket Classifier & Auto-Router | FastAPI, Next.js, PostgreSQL, MLflow, Docker, AWS

2026

- Built an end-to-end MLOps system to classify support tickets by category and urgency with confidence-based auto-routing and human-in-the-loop workflows.
- Implemented retraining pipelines with MLflow tracking and feedback-driven learning for continuous model improvement.

NexusAI: Recommendation Engine | Python, FastAPI, LangChain, Redis, PostgreSQL

2025

- Architected a recommendation API with explainable reasoning via LangChain, improving simulated user engagement metrics.
- Integrated Redis caching and a business rule dashboard to reduce latency in high-traffic environments.

Publications

Published research paper titled "A Comparative Study of Parameter-Efficient Fine-Tuning Techniques for Large Language Models" in **JETIR**. [Paper Link]

Technical Summary

AI & LLMs: Fine-tuning (QLoRA), RAG, LangChain, LangGraph, Agentic Workflows, Vertex AI
Backend & Cloud: FastAPI, Node.js, GCP (Cloud Run, Pub/Sub), Docker, CI/CD, Microservices
Databases & Systems: PostgreSQL, Pinecone, Redis, Firestore, Distributed Systems
Frontend & Languages: React, Next.js, TypeScript, Python, C++, Bash

Education

AISSMS Institute of Information Technology

July 2023 – July 2027

B.Tech in Artificial Intelligence & Data Science (CGPA: 8.05, SGPA: 8.56)

Pune, Maharashtra

Leadership & Community Engagement

Hackathons & Tech Conferences

Active Contributor

Global / India

- **Hackathons:** 1st Runner Up, Alphabyte 3.0 (PCCOE); participated in Mumbai Hacks, Google Agentic AI, Google Gen AI Exchange, and EY Techathon.
- **Open Source:** Contributed to **Google Gemini CLI** repository with PR #25313 & #26312 successfully merged.
- **Cloud Native Hackathon:** Built a production-ready microservices app with CI/CD pipelines, reducing deployment time.
- **Industry Engagement:** Attended L&T Finance event and API Conf 2025, gaining insights into multi-LLM systems.